

CS-450, CS-550, INFS-614: How to Use Oracle

1. Preliminaries

Oracle is installed on the VSE Oracle Server `apollo.vse.gmu.edu`, and is accessible from the VSE server `zeus.vse.gmu.edu`. To access Oracle you need 2 computer accounts:

1. **VSE UNIX account:** This account is created automatically for all students who are enrolled in any course offered by an VSE department. This account can be used to log into `zeus.vse.gmu.edu` using SSH. For issues regarding computer accounts see: http://labs.ite.gmu.edu/reference/faq_iteaccount.htm.
2. **Oracle account:** To become a user of the Oracle Database Management System you must [create an Oracle account](#). For problems regarding your Oracle account you may contact [Oracle administrator](#).

Note that a firewall prevents you from connecting directly to `zeus.vse.gmu.edu` from off campus. In such cases you should first connect to the GMU VPN (<http://vpn.gmu.edu/>). Then you should be able to use the `ssh` command to connect to `zeus.vse.gmu.edu`.

The VSE server `zeus.vse.gmu.edu` runs the Unix operating system, and you will need basic familiarity with this operating system. Some of the basic commands that you will want to learn are `ls`, `pwd`, `mkdir`, `cd`, `cp`, `rm`, `cat`, `more`, and `lp`, as well as one of the text editors (e.g., `emacs`, `vi`, `ed` or `pico`). If you are not familiar with any of these editors, you should consider using `pico`. Type `pico` and press return. A simple menu will appear. Once you have logged to the system, you may use the command `man` to display the manual pages for any of these commands. For example, `man man` will display the manual page of the command `man`.

2. SQL*Plus

1. **Start SQL*Plus.** At the Unix prompt, type the command `sqlplus`, and then enter your Oracle user-name and password to login to Oracle. You will get the Oracle prompt `SQL>` and you can now issue SQL commands. If you don't get the `SQL>` prompt, send email to the [Oracle Database Administrator](#).
2. **Quit SQL*Plus.** The SQL commands `quit` or `exit` will log you out of Oracle and return you to the Unix prompt. If you are done, you should now log out of `zeus.ite.gmu.edu` as well.
3. **Help.** At the `SQL>` prompt, you can type the command `help` to get on-line help messages. Type

```
SQL> help index
```

to get a list of help topics. The topics concern the `sqlplus` tool, not the SQL language.

4. **Pause output.** To instruct Oracle to pause after each page when text is displayed, assuming the size of your window is 20 lines, at the `SQL>` prompt, type

```
SQL> set pagesize 20
```

```
SQL> set pause on
```

Press return to continue after a pause, and Control-C to stop the display of text.

5. **Print.** Use the command

```
SQL> spool filename
```

to log the entire session (your input and Oracle's output) in a Unix file called `file-name.lst`. Use `spool off` to turn off the logging. Later, at the Unix prompt, you can print this file using the command

```
UNIX> enscript -P pfp file-name.lst
```

If you are working remotely, you can move this file to your local computer (using, for example, `ftp` or `sftp`) and print it there. To save paper, you may use the Unix command

```
UNIX> mpage -2 file-name.lst > file-name.ps
```

to convert your spool file to a Postscript file in which two pages are printed on a single sheet of paper, and then print with

```
UNIX> lp file-name.ps
```

If you substitute `"-2"` with `"-4"` in the above command, you will print four pages per sheet.

6. **Demo database.** To build the demo database type `demo1d` at the Unix prompt. You may then start SQL*Plus and attempt various SQL commands against this database. To remove the demo tables, type `demodrop` at the Unix

prompt.

7. **Edit your input.** Your most recent input to SQL*Plus is saved in a buffer, which you may edit and resubmit. To edit SQL*Plus buffers, use the command

```
SQL> edit
```

This command launches a text editor with the contents of the buffer. When you finish your editing, save, exit the editor, and resubmit using the command

```
SQL> run
```

The commands `edit` and `run` may be abbreviated `ed` and `r`, respectively.

The default Unix editor is `ed`. To change this default, set the `EDITOR` environment variable. For example, if you want `pico` to be your editor, at the Unix prompt type

```
UNIX> EDITOR=pico; export EDITOR
```

before you start SQL*Plus. (This assumes that you are using `bash` as your command interpreter.) You may also include the above line in your Unix configuration file `.bash_profile` so that it is executed automatically every time you login.

8. **Customize the SQL*Plus environment.** To customize your SQL*Plus environment, create a Unix file called `login.sql`. The commands in this file will be executed every time you start SQL*Plus. For example, you could include in this file the commands for pausing the output described earlier.

9. **Examples.** Below are examples of three popular operations. In each case, the commands are entered at the `SQL>` prompt.

- To create a relation, use the `create table` command. For example:

```
create table DEPT
( DEPTNO number(2),
  DNAME varchar2(14),
  LOC varchar2(13) );
```

- To insert rows into a relation, use the `insert` command. For example,

```
insert into DEPT values (20,'research','Dallas');
```

- To display the rows of a relation, use the command `select`. For example,

```
select * from DEPT;
```

You may use the `spool` command to turn on the spooling before issuing the `select` command, and then print the spool file as described earlier.

10. **Scripts.** When a sequence of SQL*Plus commands has to be repeated, script files may be used. Using a text editor, create a file of SQL*Plus commands. The file name must end with `.sql`, for example, `myscript.sql`. Then, at the Unix prompt type

```
UNIX> sqlplus @myscript
```

Or, at the SQL prompt type

```
SQL> @myscript
```

As an example, consider this script file called `table_dept.sql`:

```
drop table EMP;
create table EMP
( EMPNO number(4) not null,
  ENAME varchar2(10),
  JOB varchar2(9),
  MGR number(4),
  HIREDATE date,
  SAL number(7,2),
  COMM number(7,2),
  DEPTNO number(2) );
insert into EMP values (7369,'Smith','clerk',7902,'17-DEC-80',800,NULL,20);
insert into EMP values (7499,'Allen','salesman',7698,'20-FEB-81',1600,300,30);
```

The commands in this script file (1) delete the present `EMP` table, (2) create a new `EMP` table with the specified columns, and (3) insert two rows into the `EMP` table. To execute these commands, type

```
UNIX> sqlplus @table_dept
```

at the Unix prompt, or

```
SQL> @table_dept
```

at the SQL> prompt. This method makes it easier to initialize a database with a large number of rows. And if definitions need to be changed, it may be easier to edit the script file and re-execute it. Unless you include the SQL*Plus command quit at the end of the script, when the script is done, you will be at the SQL prompt.

11. Finding out what's in your database.

- To find out the tables in your database:

```
select TABLE_NAME
from USER_TABLES;
```

- To find out the columns in a table:

```
describe table-name;
```

- To find out the rows in a table:

```
select *
from tablename;
```

12. Circular constraints.

Suppose we want to define two tables: R(A,B) and S(C,D), in which A and C are primary keys, B is a foreign key that references S.C, and D is a foreign key that references R.A. There are two different problems.

First, it is impossible to include in the definition of R(A,B) the foreign key reference to S, because S has not been created yet. Similarly, if we attempt to create S(C,D) first, it cannot include the foreign key reference to R, because R has not been created yet. The solution is to create the first table without the foreign key constraint, and add it later, as follows:

```
create table R (
A number primary key,
B number
);

create table S (
C number primary key,
D number references R deferrable
);

alter table R
add constraint BREFC foreign key (B) references S deferrable
;
```

Note three important things:

1. We could, of course, choose to create *both* tables without the foreign keys and add both later.
2. The alter table command must include a *name* for the new constraint (in this case, we chose BREFC).
3. Both foreign key constraints were defined as *deferrable*.

Now that the tables with the circular constraints have been created, there is a second problem: how to insert rows. If we attempt to insert a row such as (1,10) into R, Oracle will complain about a foreign key violation because the table S should already have the value 10 in column C. But if we address this by inserting first a row such as (10,1) into S, Oracle will complain about a foreign key violation because the table R should already have the value 1 in column A.

This is why we defined both foreign key constraints as *deferrable*. Normally constraints are checked immediately upon each update. However, if a constraint has been defined as *deferrable*, we may request that constraint checking will be deferred. This is done with the command

```
set constraint all deferred;
```

Now we can add rows. For example:

```
insert into R values (1, 10);
insert into R values (2, 20);
insert into S values (10, 1);
insert into S values (20, 2);
```

When we are done, we should validate the content of the database by activating the constraints. This is done with the command

```
set constraint all immediate;
```

In this case the checking will succeed, and at the end of the transaction (when the SQL*Plus session ends or when we issue commit) the updates will be saved. However, if we also added

```
insert into S values (30, 3);
```

then the checking of constraints will fail. The table S may show this additional row during this session, but when the session ends (or when we issue `commit`), the transaction will be *rolled back*, and all its updates ignored!

The entire sequence of commands is repeated here:

```
create table R (
  A number primary key,
  B number
);
create table S (
  C number primary key,
  D number references R deferrable
);
alter table R
add constraint BREFC foreign key (B) references S deferrable
;
set constraint all deferred;
insert into R values (1, 10);
insert into R values (2, 20);
insert into S values (10, 1);
insert into S values (20, 2);
set constraint all immediate;
```

On a final note, when you use a script to create your database, and you include `drop table` commands at its beginning, and there are foreign key constraints, then you must drop the tables in the correct order! For example if table R has a foreign key that references table S, then you must drop table R first! However, if tables R and S are "locked" in circular constraints, then you must use the `cascade constraints` option of the `drop table` command!

3. C-Embedded SQL

Write your C-embedded SQL programs using a text editor. The program file name must end with `.pc`. To pre-compile an embedded C/SQL program `myprog.pc`, use the command

```
UNIX> occ iname=myprog.pc oname=myprog.c
```

which will generate a C program by the name `myprog.c`. To compile this program, use the command:

```
UNIX> cc -L /usr/local/oracle/OraHome1/lib -l clntsh myprog.c
```

This will generate an executable file by the name `a.out`. Finally, to execute your program, use the command

```
UNIX> ./a.out
```

Note that syntax errors in your Pro*C program may result in a confusing "Segmentation error" message.

An example C-embedded SQL program

```

/*****
This program accesses relations created by the program "demobld".
It connects to ORACLE, declares and opens a cursor, fetches the names
and salaries of all employees, displays the results, then closes the
cursor.
*****/

#include <stdio.h>

EXEC SQL BEGIN DECLARE SECTION;

    VARCHAR    username[20]; /* necessary structures for Oracle login */
    VARCHAR    password[20]; /* do not change these lines */
    VARCHAR    empname[11];  /* structures specific to this program */
    float      salary;       /* note that empname has an extra */
                           /* position for string terminator */

EXEC SQL END DECLARE SECTION;

EXEC SQL INCLUDE sqlca;      /* include SQL "communication area" for */
                           /* reporting the results of SQL execution */
                           /* after each statement */

void sqlerror();             /* handle unrecoverable errors */

main()
{
    /* Connect to ORACLE. */

    EXEC SQL WHENEVER SQLERROR DO sqlerror(); /* upon error, goto this */
                                              /* user-defined function */
                                              /* (implicit error-handling) */

    EXEC SQL CONNECT : "john" IDENTIFIED BY : "xyz"; /* substitute your name and password */

```

```
/* for "john" and "xyz" */
```

```
printf("\nConnected to ORACLE!\n");

/* begin work */

EXEC SQL DECLARE employee CURSOR FOR /* declare a query and a cursor */
SELECT ename, sal
FROM emp;

EXEC SQL OPEN employee; /* do the query (creates temp table) */

printf("\nEmployee_Name      Salary\n");
printf("-----          ----- \n");

/* Loop, fetching all employees and their salaries */

for ( ; ; )
{
    EXEC SQL WHENEVER NOT FOUND DO break; /* exit loop when no more rows */

    EXEC SQL FETCH employee INTO :empname, :salary; /* fetch a row (2 values)*/

    /* The Pro*C declaration "VARCHAR empname" created
       a structure with two fields:
       empname.arr holds the string
       empname.len holds the length of the string
       the next statement terminates the string with a null byte */

    empname.arr[empname.len] = '\0';
    printf("%-17s%9.2f\n", empname.arr, salary);
}

EXEC SQL CLOSE employee; /* delete temporary relation */

printf("\nHave a good day!\n\n");

EXEC SQL COMMIT WORK RELEASE; /* transaction succeeded */
/* commit the changes, if any, */
/* and release all locks */

exit(0);
}

void sqlerror()
{
    EXEC SQL WHENEVER SQLERROR CONTINUE;

    printf("\nORACLE error detected:\n");
    printf("\n% .70s \n", sqlca.sqlerrm.sqlerrmc);
    /* transaction failed */
    EXEC SQL ROLLBACK WORK RELEASE; /* undo the changes, if any, */
    /* and release all locks */

    exit(1);
}
```

4. Java-SQL

The JDBC package: Java provides relational database access via the package *java.sql* (this library implements JDBC, the Java DataBase Connectivity subsystem). Your Java program should include the line

```
import java.sql.*;
```

Oracle drivers: To access Oracle, you need to load Oracle-specific drivers. Your Java program should include the line

```
String driverName = "oracle.jdbc.driver.OracleDriver";
Class.forName(driverName);
```

Connecting to your database: To connect your Java program to your Oracle database, your program should include the line

```
String url = "jdbc:oracle:thin:@apollo.ite.gmu.edu:1521:ite10g";
Connection conn = DriverManager.getConnection(url,"username","password");
```

conn is a connection object that points to your database. For instance, for a student who was assigned the username "john" and the password "xyz" to connect to the oracle database, the statement would be

```
Connection conn = DriverManager.getConnection(url,"john","xyz");
```

Executing SQL statements:

- To execute SQL statements, you must first issue

```
Statement stmt = conn.createStatement();
```

stmt is a statement object that holds your future statements.

- The `executeUpdate` command is used for SQL statements that return no results (e.g., create or drop tables, insert or delete rows). For example, the command

```
stmt.executeUpdate("CREATE TABLE MYTABLE (A VARCHAR[25], B INTEGER)");
```

creates a new table called MYTABLE with columns A and B.

- The `executeQuery` command is used for SQL queries. For example, the command

```
ResultSet myresults = stmt.executeQuery("SELECT * FROM MYTABLE");
```

retrieves the content of the table MYTABLE.

Your Java-SQL program file name should end with `.java`. To compile a program called `myprog.java`, issue

```
UNIX> javac MyProg.java
```

This will create a file called `MyProg.class`. To execute your program, issue

```
UNIX> java MyProg
```

An example Java-SQL program

```
/******
```

This program accesses relations created by the program "demobld". It connects to ORACLE, fetches the names and salaries of all employees, and displays the results.

```
*****/
```

```
import java.sql.*; //Import the java SQL library
```

```
class MyProg //Create a new class to encapsulate the program
{
```

```
    public static void SQLError (Exception e) //Our function for handling SQL errors
    {
        System.out.println("ORACLE error detected:");
        e.printStackTrace();
    }
```

```
    public static void main (String args[]) //The main function
```

```
{
```

```
    try { //Keep an eye open for errors
```

```
        String driverName = "oracle.jdbc.driver.OracleDriver";
        Class.forName(driverName);
```

```
        System.out.println("Connecting to Oracle...");
```

```
        String url = "jdbc:oracle:thin:@apollo.ite.gmu.edu:1521:ite10g";
        Connection conn = DriverManager.getConnection(url,"john","xyz");
```

```
        System.out.println("Connected!");
```

```
        Statement stmt = conn.createStatement(); //Create a new statement
```

```
        //Now we execute our query and store the results in the myresults object:
        ResultSet myresults = stmt.executeQuery("SELECT ename, sal FROM emp");
```

```
        System.out.println("Employee_Name\tSalary");
        System.out.println("-----\t-----"); //Print a header
```

```
        while (myresults.next()) //pass to the next row and loop until the last
```

```
        {
            System.out.println(myresults.getString("ename") + "\t\t" + myresults.getString("sal")); //Print the current row
        }
```

```
        conn.close(); // Close our connection.
```

```
    }
    catch (Exception e) {SQLError(e);} //if any error occurred in the try..catch block, call the SQLError function
```

```
}
}
```

Last updated: Friday, 19 October 2012, 3:00pm
Amihai Motro